


Optimization

3.3.3 Conjugate gradient method (Fletcher–Reeves)

José Firmino Aguilar Madeira

The same start — where steepest descent would begin to zigzag


The example from 3.3.1: $f(x) = \frac{1}{2}(x_1^2 + 9x_2^2)$ from $x_0 = (9, 1)$, exact line search. The first step is **the same** as in steepest descent: $d_0 = -\nabla f(x_0)$, and we reach $x_1 = (7.2, -0.8)$.



At x_1 , the gradient would turn 90° and the zigzag would begin (3.3.1). The question: is there a better direction, using only gradients?

Remember the previous direction


Instead of following only $-\nabla f(x_1)$, we add a part of the previous direction: $d_1 = -\nabla f(x_1) + \beta_1 d_0$,
with $\beta_1 = \frac{\|\nabla f(x_1)\|^2}{\|\nabla f(x_0)\|^2} = \frac{103.68}{162} = 0.64$ (Fletcher-Reeves).



Extra cost: no additional gradient — only storing d_0 and $\|\nabla f(x_0)\|^2$, and two inner products.

Two iterations instead of 74


With $d_1 = (-12.96, 1.44)$ and exact line search ($\alpha_1 = 5/9$): $x_2 = (0, 0)$, the minimum, in two iterations ($n = 2$).



Steepest descent needed 74 iterations to reach $\|\nabla f\| < 10^{-6}$, with the same information per iteration.

What does “conjugate” mean?

d_0 and d_1 are not perpendicular, but they are **A-conjugate**, with $A = \nabla^2 f = \text{diag}(1, 9)$:
 $d_0^T A d_1 = (-9, -9) (-12.96, 12.96)^T = 116.64 - 116.64 = 0$. Changing the scale, $y_1 = x_1$, $y_2 = 3x_2$, the ellipses become circles — and the conjugate directions become **perpendicular**.



With circles, two perpendicular steps with exact search reach the centre. “Conjugate” = perpendicular in the geometry of A : minimizing along d_1 does not spoil what was gained along d_0 .

Now formally

Quadratic $f(x) = \frac{1}{2}x^T A x - b^T x$ with $A = A^T \succ 0$ (SPD). **A-conjugate** directions:

$$d_i^T A d_j = 0 \quad (i \neq j) \quad (\text{in the example: } d_0^T A d_1 = 0).$$

Result

SPD quadratic + exact arithmetic + exact line searches along A-conjugate directions $\Rightarrow$ **the minimum in at most n iterations.**

Fletcher and Reeves (1964) generate these directions *using only gradients* ($g_k = \nabla f(x_k)$):

$$d_0 = -g_0, \quad d_k = -g_k + \beta_k d_{k-1}, \quad \beta_k^{FR} = \frac{g_k^T g_k}{g_{k-1}^T g_{k-1}} = \frac{\|g_k\|^2}{\|g_{k-1}\|^2}.$$

Steepest descent (3.3.1): consecutive gradients orthogonal, $g_{k+1}^T g_k = 0$

Conjugate gradients: A-conjugate directions, $d_i^T A d_j = 0$: orthogonal “in the geometry of A”

The algorithm (Fletcher–Reeves)

Fletcher–Reeves: $x_0, \varepsilon_g, k_{\max}$

```
1:  $g_0 \leftarrow \nabla f(x_0); d_0 \leftarrow -g_0; k \leftarrow 0$ 
2: loop
3:    $\alpha_k \leftarrow \arg \min_{\alpha > 0} f(x_k + \alpha d_k)$   $\triangleright$  line search
4:    $x_{k+1} \leftarrow x_k + \alpha_k d_k; g_{k+1} \leftarrow \nabla f(x_{k+1})$ 
5:   if  $\|g_{k+1}\| \leq \varepsilon_g$  or  $k \geq k_{\max}$  then return  $x_{k+1}$ 
6:   end if
7:    $\beta_{k+1} \leftarrow \|g_{k+1}\|^2 / \|g_k\|^2$ 
8:    $d_{k+1} \leftarrow -g_{k+1} + \beta_{k+1} d_k; k \leftarrow k + 1$ 
9: end loop
```

$\arg \min_{\alpha > 0} \varphi(\alpha)$: the *value of α* at which φ attains its minimum — the minimizer, not the minimum value (2.1).

current gradient + **memory** of the previous direction

- Cost per iteration: that of steepest descent ($n_g = 1$ + the evaluations of f in the line search) plus two inner products.
- Memory: $x_k, g_k, d_k \sim O(n)$, no Hessian.
- On the SPD quadratic, with exact search and exact arithmetic: at most n iterations. And beyond quadratics?

Beyond quadratics: why and when to restart

Restart = set $\beta_{k+1} = 0$, that is, go back to $d_{k+1} = -g_{k+1}$. **Why:** conjugacy was built for a quadratic with exact search; on a non-quadratic f , or with inexact search, it is gradually lost and the “memory” starts to get in the way.

When to restart	Test	Why
1. Every n iterations	$(k+1) \bmod n = 0$	on a quadratic n directions suffice
2. d_{k+1} is no longer a descent direction	$g_{k+1}^T d_{k+1} \geq 0$	the line search would not decrease f
3. Consecutive gradients far from orthogonal	see note*	conjugacy has been lost



* Powell's criterion (1977): restart if $|g_{k+1}^T g_k| \geq \nu \|g_{k+1}\|^2$, with $\nu \approx 0.1$ – 0.2 (with exact search on a quadratic, $g_{k+1}^T g_k = 0$).

Effect on Rosenbrock (next page): FR **without** restart, 102 iterations; **with** restart every $n = 2, 36$. The MATLAB code in these notes uses tests 1 and 2.

Another common choice of β : Polak–Ribière⁺

$\beta_k^{PR+} = \max(0, g_k^T (g_k - g_{k-1}) / \|g_{k-1}\|^2)$ — beyond quadratics, it restarts by itself ($\beta \geq 0$): when the fraction is negative, $\beta = 0$.

Rosenbrock: steepest descent vs. conjugate gradients vs. Newton



- Rosenbrock is not quadratic, but nearly so near the minimum: FR with restart every $n = 2$ iterations converges in 36 iterations ($\|\nabla f\| < 10^{-4}$); without restart, 102; steepest descent does not converge in 3000.
- Damped Newton: 14 iterations — but each one costs a Hessian. Counted until $f < 10^{-4}$, with the same high-precision line search (Brent, tol. 10^{-10}): Newton $n_g = n_H = 12$, $n_f = 269$ in the line search (281 with f at the iterates); FR $n_g = 34$, $n_f = 684$ (718). For $n = 2$ the line search is the largest cost; as n grows, forming, storing and factorizing the Hessian makes gradient-only methods ever more attractive.

number of iterations $\neq$ cost

On the computer

The core, in MATLAB:

```
g = grad(x); d = -g;           % 1st direction
% ... at each iteration:
alpha = ...;                   % line search
x = x + alpha*d;
gn = grad(x);
beta = (gn'*gn)/(g'*g);        % Fletcher-Reeves
d = -gn + beta*d;
if mod(k+1,n)==0 || gn'*d >= 0
    d = -gn;                   % restart
end
g = gn;
```

The full function (ConjGrad, with info.nfev, info.ngcv) is in the course companion code. On the quadratic: 2 iterations, $n_g = 3$.

The line search is not a detail

On the quadratic, exact conjugacy presupposes exact line search ($g_{k+1}^T d_k = 0$). On nonlinear problems one normally uses *inexact* searches that are sufficiently accurate (Wolfe conditions), together with restarts.


With fminbnd at its default tolerance, FR on Rosenbrock may fail to converge: of the three methods studied here, it is particularly sensitive to the quality of the line search.

The methods studied on the same problem, with the cost in evaluations

Rosenbrock from $(-1.5, 2)$; stopping: first evaluation (or iterate) with $f < 10^{-4}$ (f at the iterates counts in n_f); in 3.3, accurate line search (Brent, tol. 10^{-10}). The grid (3.2.2), for exploring and initializing, is left out.

Method	Derivatives (order)	n_f	n_g	n_H	equiv. in f
Localized random [†]	0	975 [867, 1052]	—	—	975
Nelder–Mead	0	166	—	—	166
Box ($\Delta_0 = 1$)	0	13 445	—	—	13 445
Hooke–Jeeves ($P_0 = 0.5$)	0	353	—	—	353
Steepest descent [‡]	1	60 900	3000	—	72 900
Conj. grad. (FR, restart n)	1	718	34	—	854
Pure Newton	2	6	5	5	46
Damped Newton	2	281	12	12	377

[†] stochastic: median [quartiles] of 30 runs. [‡] does not reach $f < 10^{-4}$ in 3000 iterations ($f = 2.4 \cdot 10^{-4}$). Equiv. in f (central derivatives, $n = 2$): $n_f + 4n_g + 4n_H$ (the Hessian, $2n^2 + 1 = 9$ points, reuses the 4 of the gradient and $f(x_k)$).



On this problem, with these parameters: pure Newton has the lowest cost when it converges to the intended solution; Nelder–Mead beats the first-order methods because for $n = 2$ the accurate line search costs ≈ 20 evaluations per iteration. With large n ($n = 100$: numerical gradient 200, Hessian 20 001) gradient-only methods gain ground — there is no winner in the abstract.

Summary table — 3.3

What information do we have about f ? How much does it cost to obtain it? What can we guarantee with it?

Method	Information	Cost per iteration	Guarantee	Main limitation
Steepest desc. (3.3.1)	∇f	$n_g = 1$; $n_f = \text{cost of the line search; algebra } O(n)$	stationarity from any x_0 (bounded level set, Lipschitz gradient, suitable search); linear, bound $((\kappa - 1)/(\kappa + 1))^2$ (SPD quadr., exact search)	zigzag if κ large
Newton (3.3.2)	$\nabla f, \nabla^2 f$	$n_g = n_H = 1 + n \times n$ system ($O(n^3)$; memory $O(n^2)$)	local quadratic if $H(x^*) \succ 0$ (under regularity; damped: $\alpha_k \rightarrow 1$ with Armijo/Wolfe from $\alpha = 1$); 1 it. on SPD quadratics	expensive Hessian; without damping it may diverge or go to maxima/saddles
Conj. grad. (3.3.3)	∇f + previous direction	$n_g = 1$; $n_f = \text{cost of the line search; memory } O(n)$	$\leq n$ it. on SPD quadratics (exact search and arithmetic); in general, $\liminf \ \nabla f(x_k)\ = 0$ under the row-1 assumptions with exact or strong Wolfe search (Al-Baali, 1985)	sensitive to the line search

Rosenbrock, accurate line search, stop at $\|\nabla f\| < 10^{-4}$: steepest descent 3000 it. (not reached; $n_g = 3000$); damped Newton 14 it. ($n_g = 15$, $n_H = 14$); FR with restart 36 it. ($n_g = 37$). Steepest descent if n is huge and accuracy modest; Newton if the Hessian is cheap; conjugate gradients and quasi-Newton balance cost per iteration and use of curvature/memory.

KEY TAKEAWAYS — 3.3.3

1. Conjugate gradients = current gradient + **memory** of the previous direction:
$$\mathbf{d}_k = -\mathbf{g}_k + \beta_k \mathbf{d}_{k-1} \text{ (Fletcher-Reeves).}$$
2. **A-conjugate** directions: on an SPD quadratic, with exact search and exact arithmetic, the minimum in **at most n iterations**.
3. Cost per iteration almost equal to that of steepest descent, memory $O(n)$, no Hessian: an attractive compromise for large n .
4. Beyond quadratics: **restarts** and an accurate line search — of the three methods, it is the most sensitive to it.
5. Number of iterations $\neq$ cost: compare in evaluations, on this problem, with these parameters.

KEY TAKEAWAYS — Chapter 3

1. In several dimensions, in direction-based methods, each iteration chooses a **direction** and a **step**; the methods differ in the information they use.
2. **Derivative-free** (3.2): only values of f , to explore or to refine locally. NM has no general guarantee; HJ is at the origin of pattern search, whose formalized versions (GPS) have stationarity in the limit under suitable assumptions.
3. **Derivative-based** (3.3): steepest descent is cheap but slow if ill-conditioned; Newton is fast near the minimum, but expensive and needs safeguards; conjugate gradients sit in between.
4. Cost is counted in **evaluations**, not iterations; there is no winner in the abstract.
5. Local methods guarantee, at best, **stationarity** — not the global minimum.

Next (Ch. 4): **and what if not all directions are allowed?**

References Full list: *Course bibliography*, at the end of the course PDF.

- Fletcher, R., Reeves, C. M., *Function minimization by conjugate gradients*, The Computer Journal 7(2):149–154, 1964.
- Powell, M. J. D., *Restart procedures for the conjugate gradient method*, Mathematical Programming 12(1):241–254, 1977.
- Al-Baali, M., *Descent property and global convergence of the Fletcher–Reeves method with inexact line search*, IMA Journal of Numerical Analysis 5(1):121–124, 1985.
- Nocedal, J., Wright, S. J., *Numerical Optimization*, 2nd ed., Springer, 2006 — Ch. 5.

Resources

- Companion code: `ConjGrad (info.nfev, info.ngcv)`. MATLAB: `fminbnd`.
- Notebook `cap3_comparacao.ipynb`: reproduces the comparison table at the end of the chapter.